FairRecruit: Comprehensive Viva Examination Handbook

Master Technical Reference: Formulas, ML Algorithms, Internal Mechanics, Choosing Criteria, Fairness Audits & Defense

Course: Data Mining (CS) | **Project:** 05 Local Shortlisting Console & Fairness Audit Workbench | **Verified Date:** October 2026

1. Executive Summary: Overarching Goal, Accomplishments & Reality

What was the goal? In corporate recruitment, automated algorithms increasingly rank thousands of applicants. However, commercial systems frequently suffer from severe data leakage, black-box uninterpretability, and hidden algorithmic bias against protected demographic groups. Our objective was to engineer a fully local, auditable, and reproducible machine learning workbench that: (1) transparently trains and compares multiple supervised classifiers; (2) provides ranked candidate shortlists governed by explicit capacity policies; (3) performs comprehensive multi-metric fairness audits across demographic attributes; (4) extracts faithful decision paths and perturbation diagnostics; and (5) supports controlled before-and-after feature ablation experiments with version-controlled rollback.

What was achieved? We built and verified a full-stack local application (Flask REST backend + interactive vanilla JS frontend) backed by scikit-learn. Key verified achievements include: **(a) Strict Leakage Elimination:** All imputation, scaling, and one-hot encoding steps are bound into a unified pipeline fitted strictly on the 60% training partition and refitted inside every CV fold. **(b) Dual Selection Policies:** Both score thresholding (tau >= 0.50) and fixed-capacity Top-N quota ranking with deterministic tie-breaking. **(c) Three Complementary Fairness Gaps:** Rigorous mathematical computation of Selection Rate Gap (Demographic Parity), Recall Gap (Equal Opportunity), and Precision Gap (Predictive Parity), alongside 95% Wilson confidence intervals. **(d) Isolated Ablation Engine:** Controlled retraining without specific predictors (e.g. city development index) with active manifest versioning and SQLite audit logging. **(e) Unlabeled Candidate Handling:** Uploaded applicant pools are evaluated for selection without fabricating negative labels or corrupting the trained model.

CRITICAL VIVA DISTINCTION: The Surrogate Label Limitation

The underlying demonstration data is the Kaggle HR Analytics dataset (19,158 records). In this dataset, target = 1 denotes an enrollee **looking for a job change**, while target = 0 denotes someone not seeking a change. In our console, this is treated as a surrogate for candidate shortlisting. **Viva Defense Rule:** Never tell the examiner this model measures candidate qualification, technical competence, or job performance. Real hiring systems require validated competency scores, defensible job descriptions, and legal compliance. Conflating job-seeking interest with hiring suitability is a fatal conceptual flaw.

2. Data Architecture, Preprocessing & Leakage Elimination

The dataset consists of **19,158 tabular records**. Data preprocessing enforces strict chronological hygiene. The data is partitioned using a stratified **60% Training (11,494), 20% Validation (3,832), and 20% Test (3,832)** split (Random State: 42). From 10 active model predictors, preprocessing generates 45 encoded features.

Variable Name	Raw Format	Preprocessing & Cleaning Technique	Role in Console
enrollee_id	Integer ID	Preserved for tracking and tie-breaking; strictly excluded from feature matrices.	Excluded / Identifier
gender	String ('Male', 'Female', 'Other')	Missing values (~23.5%) mapped to explicit 'Unknown'. Retained for audits.	Excluded from Inputs (Audit Only)
city_development_index	Float (0.450 - 0.949)	Continuous urban development score; median imputed if missing.	Active Input Predictor
experience	Censored String ('<1', '>20')	Censored strings mapped: '<1' -> 0, '>20' -> 21. Numeric years.	Active Input Predictor
training_hours	Integer (1 - 336)	Continuous completed hours; median imputed. Outlier IQR fence flagged.	Active Input Predictor
education_level	Categorical (5 levels)	Missing mapped to 'Unknown'; One-Hot Encoded (handle_unknown='ignore').	Active Input (One-Hot)
company_size / type	Categorical (8 / 6 levels)	Missing mapped to 'Unknown'; One-Hot Encoded with 0/1 dummy columns.	Active Input (One-Hot)
relevent_experience	Binary Categorical	Mapped to dummy binary indicator ('Has relevant experience').	Active Input (One-Hot)

Key Data Engineering Techniques Explained:

- **Explicit Missing Token ('Unknown'):** 4,508 records (23.53%) lack gender information. Rather than deleting these rows (which would introduce selection bias) or guessing values, they are mapped to an explicit 'Unknown' token. This allows auditing selection disparity among candidates who decline to state demographics.
- **Boundary Mapping for Censored Categories:** Observations such as '<1' and '>20' are censored ranges, not erroneous data. Mapping them to numeric boundaries (0 and 21) preserves monotonic ordering while documenting exact provenance in _prov_ columns.
- **Outlier Detection via IQR Fences:** For skewed features (training_hours), outlier fences are computed: [Q1 - 1.5*IQR, Q3 + 1.5*IQR]. Extreme values are flagged for human inspection, but not arbitrarily deleted. *Examiner note:* Standard scaling does NOT remove outliers; it merely rescales the distribution.
- **Elimination of Data Leakage:** In naive machine learning implementations, scalers and imputers are fit across the entire dataset prior to splitting. This constitutes data leakage because validation and test distributions contaminate the training process. In our revised architecture, SimpleImputer, StandardScaler, and OneHotEncoder are fitted strictly on the 60% training split via a scikit-learn Pipeline. Test and live upload data only undergo transform(). Inside cross-validation, the entire pipeline is refitted independently inside every fold.

3. Machine Learning Techniques: What Happens Inside Each Algorithm

The platform implements five distinct supervised classification algorithms representing different inductive paradigms. Here is the detailed mathematical and algorithmic breakdown of what happens inside each technique:

A. Decision Tree Classifier (Active Default Model)

Algorithmic Workflow: The decision tree builds a top-down binary tree by recursively splitting the training records. At every internal node, the algorithm iterates over all 45 encoded features and all possible split thresholds. It evaluates the weighted reduction in Gini Impurity and selects the exact feature and cutoff that maximizes this reduction. Splitting terminates when maximum depth (5) is reached or when a split would produce a leaf with fewer than 25 samples.

Decision Tree: Gini Impurity & Splitting Mathematics

Gini Impurity Formula & Weighted Split Criterion:

I_G(Node) = 1 - SUM_{c=0,1} (p_c)^2 where p_c is the proportion of class c in that node.

- Pure node (100% class 0 or 100% class 1): I_G = 1 - (1.0^2 + 0^2) = 0.0 (minimum impurity).
- Maximum impurity (50% class 0, 50% class 1): I_G = 1 - (0.5^2 + 0.5^2) = 0.50.

Weighted Impurity Reduction (Information Gain):

Delta I_G = I_G(Parent) - [(N_Left / N_Parent) * I_G(Left) + (N_Right / N_Parent) * I_G(Right)]

The tree evaluates every feature cutoff and chooses the split that maximizes Delta I_G.

Worked Manual Calculation: Gini Split Reduction

Step-by-Step Hand-Worked Calculation of Gini Impurity Reduction (Viva Favorite):

Suppose a parent node has **100 applicants: 25 positive (target=1) and 75 negative (target=0)**.

1. Parent Impurity: I_G(Parent) = 1 - [(25/100)^2 + (75/100)^2] = 1 - [0.0625 + 0.5625] = **0.3750**.
2. A split on 'city_development_index <= 0.625' divides the node into:
   - **Left Child (40 applicants):** 20 positive, 20 negative.
     I_G(Left) = 1 - [(20/40)^2 + (20/40)^2] = 1 - [0.25 + 0.25] = **0.5000**.
   - **Right Child (60 applicants):** 5 positive, 55 negative.
     I_G(Right) = 1 - [(5/60)^2 + (55/60)^2] = 1 - [0.00694 + 0.84028] = **0.1528**.
3. Weighted Child Impurity: (40/100) * 0.5000 + (60/100) * 0.1528 = 0.2000 + 0.0917 = **0.2917**.
4. Impurity Reduction: Delta I_G = 0.3750 - 0.2917 = **+0.0833**. The tree commits this split.

B. Logistic Regression (Linear Log-Odds & Sigmoid)

Algorithmic Workflow: Models the log-odds of the positive class as a linear function of features: **ln(p / (1 - p)) = w^T * x + b**. Inverting this yields the standard logistic sigmoid: **p = 1 / (1 + e^(-(w^T * x + b)))**. Parameters are optimized using L-BFGS to minimize L2-regularized Binary Cross-Entropy (Log Loss): **Loss = - (1/N) * sum [y_i * ln(p_i) + (1 - y_i) * ln(1 - p_i)] + (lambda / 2) * ||w||^2**. *Why Scaling is Vital:* Feature weights reflect gradient magnitudes. Without StandardScaler, features with large numeric ranges (training_hours: 1-336) would dominate the penalty term relative to compact features (city_index: 0.45-0.95), corrupting the learned coefficients.

C. Random Forest (Bootstrap Bagging Ensemble)

Algorithmic Workflow: Constructs an ensemble of 80 decorrelated decision trees. Each tree is trained on a bootstrap sample (random sampling with replacement of size N from training data). At each node split, the tree is constrained to consider only a random subset of features (m = sqrt(45) approx 7 features). Predictions are formed by averaging probability estimates across all 80 trees: **p(x) = (1/80) * sum_{t=1..80} p_t(x)**. *Tradeoff:* Reduces variance and boosts AUC (0.7845), but sacrifices exact decision-path interpretability, functioning as a black box.

D. Gaussian Naive Bayes (Conditional Independence)

Algorithmic Workflow: Applies Bayes' Theorem under the strict assumption that all features are conditionally independent given the target: **P(Y=1 | x) = [P(Y=1) * prod P(x_i | Y=1)] / P(x)**. Continuous feature likelihoods are modeled as univariate normal distributions: **P(x_i | y) = (1 / sqrt(2*pi*sigma_y^2)) * exp(- (x_i - mu_y)^2 / (2*sigma_y^2))**.
Variance Smoothing (var_smoothing = 1e-3): Adds 10^-3 of the maximum feature variance to all variances, preventing zero-division when class variance is near zero. *Limitation:* Real-world one-hot dummy variables violate independence, resulting in poor probability calibration and high false-positive rates.

E. K-Nearest Neighbors (KNN - Instance-Based Learning)

Algorithmic Workflow: Non-parametric algorithm that memorizes scaled training instances. To score an applicant, it calculates the Euclidean distance **d(x, x') = sqrt(sum (x_j - x'_j)^2)** to all 11,494 training records. The k=7 nearest neighbors vote with inverse-distance weighting: **w_i = 1 / d(x, x_i)**. **Bug Fixed:** The original repository standardized train, test, and live applicants separately (a single row standardized against itself produced all zeros!). Our unified pipeline fixes this, guaranteeing exact consistency between batch and single-applicant predictions.

F. Dimensionality Reduction: PCA vs. Train-Fitted Feature Selection

We evaluate two dimensionality reduction methods targeting k=10 dimensions on the training partition:
• **Train-Fitted Importance Selection (SelectFromModel):** Fits a Decision Tree and extracts Gini feature importances. It retains the top 10 most predictive original features (e.g. city_development_index, experience, company_size). Features preserve their natural human names.
• **Principal Component Analysis (PCA):** Computes the covariance matrix of scaled training data, solves the eigenvalue equation **C * v = lambda * v**, and projects data onto the 10 orthogonal eigenvectors of greatest variance (explaining ~58.4% of variance). *Examiner Note:* While PCA minimizes reconstruction error, every principal component is a linear mixture of all 45 original variables, destroying direct interpretability.

4. Performance Evaluation Metrics & Exact Mathematical Formulas

All candidate predictions are evaluated using confusion matrix counts: True Positives (TP), False Positives (FP), False Negatives (FN), and True Negatives (TN).

Metric Name	Mathematical Formula	Recruitment Interpretation	Edge Case & Handling
Accuracy	(TP + TN) / Total	Overall proportion of correct candidate classifications.	Deceptive under imbalance (Majority Baseline Trap).
Precision (PPV)	TP / (TP + FP)	Of those shortlisted, what fraction are actual positive matches? Measures interview waste.	Undefined if TP + FP = 0. Output 'Unavailable', never 0.0.
Recall (TPR)	TP / (TP + FN)	Of all genuine positive matches, what fraction did we shortlist? Measures talent loss.	Undefined if TP + FN = 0. Output 'Unavailable'.
Specificity (TNR)	TN / (TN + FP)	Proportion of negative candidates correctly filtered out.	Measures resistance to false alarms.
F1-Score	2 * (Precision * Recall) / (P + R)	Harmonic mean balancing precision and recall.	Undefined if P + R = 0.
ROC-AUC	P(Score(Pos) > Score(Neg)) + 0.5*P(Tie)	Threshold-independent ranking discrimination probability.	Requires both positive and negative held-out cases.

Step-by-Step Manual Calculation Walkthrough

Worked Hand-Calculation Example (Must Rehearse for Viva):
Suppose Demographic Group A has 4 applicants, 3 shortlisted, 2 True Positives, and 2 actual positives in ground truth.
Suppose Demographic Group B has 4 applicants, 1 shortlisted, 1 True Positive, and 2 actual positives in ground truth.
• **Group A Selection Rate:** (TP + FP) / N = 3 / 4 = **75.0%**.
• **Group A Recall:** TP / (TP + FN) = 2 / 2 = **100.0%**. **Group A Precision:** TP / (TP + FP) = 2 / 3 = **66.67%**.
• **Group B Selection Rate:** (TP + FP) / N = 1 / 4 = **25.0%**.
• **Group B Recall:** TP / (TP + FN) = 1 / 2 = **50.0%**. **Group B Precision:** TP / (TP + FP) = 1 / 1 = **100.0%**.
• **Fairness Gaps:** Selection Rate Gap = 75% - 25% = **50.0 pp**. Recall Gap = 100% - 50% = **50.0 pp**. Precision Gap = 100% - 66.67% = **33.33 pp**.
• **Four-Fifths Ratio for Group B:** 25% / 75% = **0.3333** (Triggers Adverse Impact screening flag because 0.3333 < 0.80).

The Majority Baseline Trap: In our training split, 75.07% of candidates have target = 0 and 24.93% have target = 1. A trivial dummy model that rejects every applicant achieves **75.07% accuracy** while finding zero talent (Recall = 0.0%). Therefore, our Decision Tree's 77.53% validation accuracy is only meaningful when paired with its positive-class Recall (55.65%), Precision (54.90%), and ROC-AUC (0.7760).

5. Choosing Criteria: Model Selection, Thresholds & Quota Policies

How did we end up with our choosing criteria? Model selection and decision rules are strictly decoupled to allow transparent tradeoffs.

Model Family	Val Accuracy	Val Precision	Val Recall	Val ROC-AUC	5-Fold CV Accuracy	Choosing Assessment
Decision Tree (Active)	77.53%	54.90%	55.65%	0.7760	78.24% +/- 0.77%	CHOSEN DEFAULT: Best balance of AUC, balanced recall, and direct explainability.
Logistic Regression	76.88%	53.20%	49.12%	0.7610	76.45% +/- 0.62%	High precision, but linear assumption misses complex feature interactions.
Random Forest	78.10%	57.40%	51.80%	0.7845	78.60% +/- 0.55%	Higher AUC, but ensemble structure acts as a black box; explanations are heuristic.
Gaussian Naive Bayes	71.20%	42.15%	68.40%	0.7320	71.50% +/- 1.10%	High recall at cost of severe false positives; violates independence assumption.
K-Nearest Neighbors	74.80%	48.90%	47.20%	0.7250	74.30% +/- 0.85%	Slow inference latency O(N*d); sensitive to local density variations.

Selection Policies: Threshold vs. Top-N Capacity Quota:

Recruitment workflows encounter two distinct operating environments, formalized in SelectionPolicy (backend/policy.py):
• **1. Probability Threshold Policy (Default: tau = 0.50):** Shortlists candidate if P(Y = 1 | x) >= tau. This evaluates each applicant independently against an absolute qualification standard. Lowering the threshold (e.g. to 0.35) increases selection count and recall, but admits more false positives, reducing precision. The ROC curve maps this exact tradeoff.
• **2. Top-N Capacity Quota Policy:** In real hiring, an interview committee has fixed operational capacity (e.g. exactly 50 interview slots). The Top-N policy sorts all candidates by predicted positive probability descending and selects the top N. **Top-N overrides thresholding.**
• **Deterministic Tie-Breaking Rule:** When two candidates have identical probabilities, ties are broken first by unrounded float score descending, then by enrollee_id ascending (lexicographical string sort). This ensures 100% deterministic, auditable decisions.
• **Single Applicant Limitation:** A single applicant cannot receive a Top-N shortlisting verdict in isolation because their rank depends on the scores of the entire candidate pool.

Cost-Optimal Threshold Analysis & Its Operational Limits

Cost-Optimal Threshold Formulation & The 0.06 Trap:

Expected Cost = FP * C_FP + FN * C_FN - TP * B_TP + TN * C_TN

Using standard assumed recruitment parameters: C_FP = $500 (wasted interview cost), C_FN = $3,500 (cost of missing a strong hire), B_TP = $5,000 (net benefit of successful hire), C_TN = $0.

Grid Search & The 0.06 Trap: The console evaluates expected cost across 101 threshold steps (0.00 to 1.00). *Examiner Trap:* Without capacity constraints, because B_TP ($5,000) and C_FN ($3,500) heavily penalize misses, the mathematically 'optimal' threshold drops to **0.06**, which shortlists 100% of candidates! This proves that cost-optimal thresholding is only operationally valid when bounded by maximum interview capacity constraints.

6. Algorithmic Fairness Audits, Disparity Gaps & Uncertainty

Auditing fairness requires measuring decisions across protected demographic slices (Gender, derived Age Band, Education Institution). Our system computes three distinct, non-congruent fairness criteria:

Fairness Criterion	Mathematical Definition	What It Measures	Recruitment Implication
Demographic Parity (Selection Rate Gap)	Max(SR_g) - Min(SR_g) where SR_g = (TP_g + FP_g) / N_g	Checks whether groups are selected at identical rates regardless of qualification.	Focuses on equal outcome distribution; ignores underlying label differences.
Equal Opportunity (Recall Gap)	Max(Recall_g) - Min(Recall_g) where Recall_g = TP_g / (TP_g + FN_g)	Checks whether qualified candidates across groups have equal chances of selection.	Protects qualified applicants from disparate false negative rates.
Predictive Parity (Precision Gap)	Max(Prec_g) - Min(Prec_g) where Prec_g = TP_g / (TP_g + FP_g)	Checks whether shortlisted candidates across groups have equal chances of success.	Ensures interview slots are equally efficient across demographic groups.
Four-Fifths Rule (Adverse Impact Ratio)	Ratio_g = SR_g / Max(SR_all) Flag triggered if Ratio < 0.80	EEOC heuristic screening for adverse selection rate impact.	Screening indicator only; NOT legal proof of unlawful discrimination.

Mathematical Impossibility Theorem of Fairness:

A foundational result in fairness theory (Chouldechova 2017, Kleinberg et al. 2016) proves that **Demographic Parity, Equal Opportunity, and Predictive Parity are mutually incompatible** whenever demographic groups have different base prevalence rates (P(Y = 1 | Group A) != P(Y = 1 | Group B)) and classification is imperfect. A model that achieves equal selection rates will inevitably suffer recall or precision gaps if one group has higher ground-truth suitability. Hiring teams must explicitly prioritize which fairness criterion aligns with policy goals rather than claiming 'absolute fairness.'

Statistical Rigor: Wilson Intervals for Denominator Uncertainty

Statistical Uncertainty & The Wilson Score Interval:
Small subgroups produce highly volatile metric estimates. A group with N=39 total members and only 4 selections has an observed precision of 100% (4/4), yet its statistical uncertainty is enormous. To prevent misleading conclusions, the console computes 95% Wilson Score Confidence Intervals:

p_mid = [p + (z^2 / (2n))] / [1 + (z^2 / n)]

Margin = [z / (1 + z^2/n)] * sqrt([p(1 - p) / n] + [z^2 / (4n^2)])

where z = 1.95996 for 95% coverage. Furthermore, 1,000-iteration plug-in binomial bootstrap intervals are computed for all metric gaps.

7. Fairness Interventions, Ablation Experiments & Influencing Factors

What factors influence fairness and predictive performance? Three primary mechanisms drive disparate outcomes:
1. **Proxy Correlations:** Unprotected features (e.g. city_development_index) often correlate strongly with protected attributes or historical advantages.
2. **Base Rate Disparities:** Differences in historical job-change interest across demographic groups create natural mathematical tension between parity and precision.
3. **Selection Thresholds:** Stricter thresholds (e.g. 0.70) amplify disparities because minority representations in the extreme score tail may be disproportionately low.

The Central Ablation Experiment: Dropping City Development Index

To evaluate mitigation, we perform an isolated ablation: retraining the Decision Tree on the identical 60% training split after removing **city_development_index** (the highest-ranking predictor), evaluated on the identical 20% validation partition at the same 0.50 threshold.

Metric / Audit Dimension	Before Ablation (Baseline)	After Dropping CDI	Observed Delta	Technical Interpretation
Validation Accuracy	77.53%	76.25%	-1.28 pp	Predictive accuracy moderately degraded.
Validation ROC-AUC	0.7760	0.7086	-0.0674	Ranking discrimination significantly impaired.
Gender Selection Rate Gap	17.62 pp	7.74 pp	-9.88 pp	DEMOGRAPHIC PARITY IMPROVED! Selection rates became much more equal.
Gender Recall Gap	20.15 pp	24.58 pp	+4.43 pp	EQUAL OPPORTUNITY WORSENED! Qualified candidates face greater disparity.
Gender Precision Gap	12.88 pp	15.52 pp	+2.64 pp	PREDICTIVE PARITY WORSENED! Interview quality disparity increased.
Four-Fifths Minimum Ratio	0.5226	0.6840	+0.1614	Adverse impact ratio improved, but remains below the 0.80 legal screening line.

THE CENTRAL VIVA LESSON: Fairness Through Omission is a Myth!

Dropping a sensitive or proxy attribute is NOT an automatic fairness solution. While removing city_development_index reduced the selection rate gap by 9.88 percentage points, it degraded overall ROC-AUC by 0.067 and **worsened both the recall gap (+4.43 pp) and precision gap (+2.64 pp)**. This proves that blind omission can harm qualified applicants while degrading predictive utility. Interventions represent multi-objective tradeoffs.

Proxy Detection via Normalized Mutual Information (NMI):

When removing an attribute, should correlated proxies also be removed? The console computes Normalized Mutual Information between the candidate column and all other active predictors on training data (using numeric quintile binning). **NMI(X; Y) = 2 * I(X; Y) / [H(X) + H(Y)]**. Attributes exceeding NMI >= 0.30 are flagged as proxy candidates. *Examiner Note:* Automatic proxy removal should be approached cautiously: dropping associated features indiscriminately strips away valid predictive signal and destabilizes model calibration.

8. Explainability: Decision Paths & Perturbation Diagnostics

The console rejects fabricated or hardcoded explanations in favor of two mathematically faithful mechanisms:

• **1. Exact Decision Tree Paths:** For any shortlisted applicant, the console traverses the actual fitted scikit-learn decision tree using decision_path(). It extracts the exact sequence of boolean inequalities that directed the applicant from the root node to their terminal leaf. For example: *'City development index > 0.625' -> 'Years of experience > 4.5' -> 'Education level is Graduate'*. This provides 100% faithful verification of how the model arrived at the leaf's class probability.

• **2. Source-Column Perturbation Diagnostics:** To measure the sensitivity of an applicant's score to a specific attribute, the console creates a perturbed copy where the target attribute is replaced by its training reference value (median for numeric, mode for categorical). The diagnostic score is the measured marginal output delta: **Delta P = P(original) - P(perturbed)**. If replacing experience with the training median drops score by 0.14, experience contributed a +0.14 positive lift for this candidate.
Examiner Caution: Perturbation measures model behavior, NOT real-world causal attribution. Replacing a single input does not simulate real-world interventions.

9. Master Viva Q&A; Guide: Top 15 Examiner Questions & Model Answers

Q1: What exactly is the data mining task, and how is ranking performed?

The core mining task is supervised binary classification on tabular records. The console fits a probabilistic classifier to estimate P(target = 1 | x). Ranking is achieved by sorting applicants in descending order of their predicted positive probabilities. It is not a separately trained learning-to-rank algorithm (such as LambdaMART).

Q2: What does target = 1 actually represent, and why is this distinction critical?

In the Kaggle dataset, target = 1 denotes an applicant looking for a job change. In our workbench, it serves as a course demonstration surrogate for candidate shortlisting. It does not represent validated job suitability or interview success. Conflating job-seeking interest with candidate qualification is a major ethical and operational failure.

Q3: Why was the Decision Tree chosen as the primary model over Random Forest or XGBoost?

Decision Tree achieved competitive validation ROC-AUC (0.7760) and balanced recall (55.65%) while offering complete white-box interpretability. In hiring audits, stakeholders must inspect exact boolean decision paths. Random Forest yielded marginally higher AUC (0.7845) but operates as an ensemble black-box where explanations are purely heuristic.

Q4: How does your implementation prevent data leakage across partitions?

We strictly partition data into 60% Train, 20% Validation, and 20% Test before any transformations occur. All preprocessing steps (median imputation, standard scaling, one-hot encoding) are fitted strictly on the training partition within a unified scikit-learn Pipeline. Inside cross-validation, the pipeline is refitted locally within each training fold.

Q5: How does the Gini Impurity split criterion work mathematically in the Decision Tree?

Gini Impurity measures class heterogeneity: I_G = 1 - sum(p_i^2). At each node, the tree evaluates all candidate feature thresholds and selects the split that maximizes weighted impurity reduction: Delta I_G = I_G(parent) - [(N_left/N)*I_G(left) + (N_right/N)*I_G(right)].

Q6: Why is K-Nearest Neighbors sensitive to feature scaling, and what bug was fixed?

KNN computes distances in Euclidean space: d = sqrt(sum (x_j - x'_j)^2). Features with large ranges (e.g. training hours up to 336) overwhelm features with small ranges (city index 0.45-0.95). The original codebase standardized train, test, and single records separately (single records became all zeros!). We fixed this by binding StandardScaler into the train-fitted Pipeline.

Q7: Why is reporting 77.53% accuracy alone deceptive?

Because of the Majority Class Baseline: 75.07% of dataset candidates are negative (target = 0). A naive classifier that rejects everyone achieves 75.07% accuracy while finding zero talent. Our model's improvement (+2.46 pp) is only defensible when paired with positive-class Recall (55.65%), Precision (54.90%), and ROC-AUC (0.7760).

Q8: How does the Top-N Capacity Quota policy differ from the probability threshold?

A threshold policy (e.g. >= 0.50) evaluates applicants independently against an absolute bar, producing variable shortlist sizes. A Top-N policy selects the top N highest-scoring candidates to match fixed interview capacity. Top-N overrides thresholds. Ties are broken deterministically by unrounded probability descending, then applicant ID ascending.

Q9: What happens when the selection threshold is lowered from 0.50 to 0.35?

The shortlisted set expands monotonically: Recall and Selection Rate cannot decrease. However, Precision typically degrades as more False Positives enter the pool. Fairness gaps can shift in either direction depending on the demographic score distributions. ROC-AUC remains unchanged because relative score ranking is unaffected.

Q10: Why did dropping city_development_index improve Demographic Parity but worsen Equal Opportunity?

City development index was strongly correlated with both target and demographic groups. Removing it compressed the score distribution between groups, reducing the Selection Rate Gap from 17.62 pp to 7.74 pp. However, because it was a legitimate predictor of positive outcomes, removing it caused more misclassifications among qualified candidates, increasing the Recall Gap from 20.15 pp to 24.58 pp.

Q11: Why is an overall sample size of N >= 30 insufficient to prove subgroup reliability?

Recall depends on actual positives (TP + FN); Precision depends on selected candidates (TP + FP). A subgroup with N=39 total members but only 4 shortlisted candidates has an effective precision denominator of only 4! Calculating precision on 4 instances produces massive sampling error, which is why the console displays 95% Wilson Score Confidence Intervals.

Q12: Why must missing gender be audited as an explicit 'Unknown' category rather than dropped?

Missing gender comprises 23.53% of the dataset (4,508 records). In our audit, the Unknown group exhibited the highest selection rate. Dropping Unknown records would distort reference baselines, artificially deflate measured maximum disparity gaps, and hide biased shortlisting among applicants who withhold demographic information.

Q13: What does the Four-Fifths Rule (Adverse Impact Ratio) prove legally and statistically?

Under US EEOC guidelines, if a group's selection rate is less than 80% (4/5ths) of the highest group's rate, it triggers an adverse impact screening flag. Statistically and operationally, it is an exploratory screening heuristic that demands investigation; it is NOT legal proof of unlawful discrimination, nor does a ratio > 0.80 prove fairness.

Q14: What is the mathematical difference between PCA and Feature Selection?

Feature Selection (SelectFromModel) ranks original encoded features by Gini importance and discards low-scoring columns, preserving exact feature semantics. PCA computes eigenvectors of the covariance matrix and projects data into orthogonal linear combinations that maximize variance. PCA destroys individual feature interpretability in hiring decisions.

Q15: How does the console handle live, unlabeled applicant uploads?

Uploaded applicant CSVs lack target labels. The console parses them through the fitted preprocessing pipeline and computes scores and Top-N rankings. Selection rates and demographic parity remain fully auditable. However, Recall and Precision are displayed as 'Unavailable' because label-dependent metrics cannot be calculated without ground truth.

10. Common Viva Pitfalls: Claims to NEVER Make

Examiners frequently probe candidates for over-claims or conceptual misunderstandings. Avoid these fatal errors:

Fatal Claim to Avoid in Viva	Why the Claim is Mathematically or Factually False	Defensible Statement to Use Instead
'Our model predicts candidate job suitability and hiring success.'	The Kaggle target measures whether an enrollee is seeking a job change, not candidate competence, interview performance, or job retention.	'The model demonstrates binary classification on a job-change surrogate dataset under an educational recruitment framing.'
'Excluding gender from predictors guarantees the model is fair.'	Fairness through Omission fails because unprotected features (city index, experience) act as proxies, preserving historical disparities.	'Audit-only protected attributes prevent direct disparate treatment, but disparate impact and proxy disparities must still be actively audited.'
'A Four-Fifths ratio above 0.80 proves our shortlist is completely fair.'	The 80% rule is an administrative screening heuristic; it ignores recall disparities, precision gaps, and individual fairness.	'The 4/5ths ratio is one exploratory screening check alongside multi-metric gap audits and Wilson confidence intervals.'
'Our 5-fold cross-validation standard deviation is a confidence interval.'	Fold standard deviation measures variability across training partitions; it is not a parametric confidence interval for deployment error.	'The +/- 0.77% spread reflects observed standard deviation across 5 stratified training folds.'
'N >= 30 means all our group fairness metrics are statistically reliable.'	Precision denominator is shortlisted count (TP+FP); Recall denominator is positive count (TP+FN). These can be tiny even when total N > 30.	'We inspect metric-specific denominators and display 95% Wilson score intervals to convey small-sample uncertainty.'
'Our perturbation diagnostics prove which features caused the hiring decision.'	Perturbations measure model output sensitivity under synthetic marginal feature replacement; they are NOT real-world causal proofs.	'Perturbation diagnostics quantify model sensitivity to specific inputs, but do not establish real-world causal mechanisms.'

11. Recommended 7-Minute Viva Presentation Script & Sequence

Structure your live project defense into this proven 7-minute cadence:

- **Minute 0:00 - 1:00 (Business Problem & Surrogate Label):** State the project scope: a local, transparent shortlisting console and fairness workbench. Explicitly state that the bundled dataset measures job-change interest as an educational surrogate for hiring decisions.
- **Minute 1:00 - 2:00 (Data Engineering & Preprocessing):** Show the 60/20/20 split (11,494 / 3,832 / 3,832). Explain the handling of censored experience ('<1', '>20'), the 23.5% missing gender retained as 'Unknown', and how unified train-fitted pipelines prevent data leakage.
- **Minute 2:00 - 3:00 (Model Benchmark & Decision Tree Selection):** Review the comparison table. Explain why the Decision Tree was selected (ROC-AUC 0.7760, accuracy 77.53% vs 75.07% majority baseline, depth 5, min leaf 25, transparent boolean rules).
- **Minute 3:00 - 4:00 (Shortlisting Policies & Tie-Breaking):** Demonstrate the candidate ranking table. Contrast thresholding (>= 0.50) with Top-N capacity quota shortlisting. Explain the deterministic tie-breaking rule.
- **Minute 4:00 - 5:00 (Fairness Auditing & Three Gaps):** Open the fairness audit view. Explain the three gaps: Selection Rate Gap (17.62 pp), Recall Gap (20.15 pp), and Precision Gap (12.88 pp). Point out the Wilson confidence intervals and the Unknown group.
- **Minute 5:00 - 6:00 (The Central Ablation Experiment):** Run the CDI ablation preview. Highlight the central finding: selection gap dropped to 7.74 pp, but recall gap worsened to 24.58 pp and AUC fell to 0.7086. Explain why 'fairness through omission is a myth.'

Show versioned apply/revert.

- **Minute 6:00 - 7:00 (Explainability & Live Uploads):** Open a shortlisted applicant drawer. Walk through their exact decision tree path and perturbation sensitivity. Conclude by demonstrating an unlabeled candidate upload where selection rates work but label metrics are marked unavailable.

12. Architecture & Codebase Dataflow Guide

Trace the execution flow across backend modules for complete technical command during code defense:

Backend Module	Primary Class / Functions	Architectural Role & Guarantees
backend/data_manager.py	DataManager, clean_row, AUDIT_LABELS	CSV loading, schema validation, boundary mapping ('<1'->0, '>20'->21), deterministic 60/20/20 split.
backend/ml_engine.py	ModelEngine, make_pipeline, evaluate, run_kfold_cv	Scikit-learn Pipeline construction, median imputation, scaling, 5-fold stratified CV, SelectFromModel, PCA.
backend/policy.py	SelectionPolicy (select, parse, to_dict)	Centralized decision logic. Enforces threshold vs Top-N quota and deterministic tie-breaking.
backend/fairness_audit.py	audit_single_attribute, wilson, compute_before_after_ablation	Per-group rates, 95% Wilson intervals, 1000-draw bootstrap gap intervals, driver diagnostics, isolated ablation.
backend/explainability.py	explain_row, decision_path traversal	Exact boolean inequality paths from fitted tree; single-attribute marginal perturbation deltas against median/mode.
backend/storage.py	SQLite persistence, model versioning manifest	Tracks active model UUID, parameter history, audit logs, and supports one-click rollback of ablated models.
app.py	Flask REST API routing	Exposes /api/train, /api/predict, /api/shortlist, /api/fairness, /api/ablation, /api/candidate//explain.

13. Quick-Lookup Master Summary & Final Viva Checklist

Verified Quantity	Exact Value	Recruitment & Evaluation Context
Total Raw Records	19,158	Full dataset size from Kaggle HR Analytics benchmark.
Stratified Split (60 / 20 / 20)	11,494 / 3,832 / 3,832	Deterministic split (Seed 42); prevents test data snooping.
Base Rate (Pos / Neg)	24.93% / 75.07%	Majority class baseline is 75.07%. Always report baseline delta.
Active Model Configuration	Decision Tree (depth=5, leaf=25)	Gini splitting; max 32 leaves; transparent boolean rule extraction.
Validation Performance	Acc: 77.53% | ROC-AUC: 0.7760	Precision: 54.90% | Recall: 55.65% (balanced operational performance).
Test Final Performance	Acc: 79.36% | ROC-AUC: 0.8100	Precision: 58.49% | Recall: 59.16% (evaluated once on held-out test split).
5-Fold Stratified CV Spread	78.24% +/- 0.77%	Observed fold SD across training folds; NOT a confidence interval.
Missing Gender Records	4,508 records (23.53%)	Audited explicitly as 'Unknown'; had highest group selection rate.
CDI Ablation Delta (Gender)	Gap: 17.62 pp -> 7.74 pp (-9.88 pp)	Selection parity improved, but Recall gap worsened from 20.15 pp to 24.58 pp.

Viva Defense Final Memory Checklist

Final 60-Second Viva Room Checklist:
[x] **Surrogate target:** Job-change interest != verified job qualification.
[x] **Gini Impurity:** I_G = 1 - sum(p_i^2); calculate parent vs weighted child split reduction by hand.
[x] **Accuracy alone:** 75.07% majority baseline means predicting all-zeros gets 75.07% accuracy with 0% recall.
[x] **Four-fifths rule:** 80% adverse impact ratio is an administrative screening flag, NOT legal proof of discrimination.
[x] **Fairness through omission:** Dropping city index improved selection gap by 9.88 pp but worsened recall gap by 4.43 pp.
[x] **Decision policy:** Top-N overrides threshold; tie-breaking uses score descending, then enrollee_id ascending.